

Faizan Munsaf

+92 332 8893064 | faizanmunsaf@gmail.com

[LinkedIn](#) | [GitHub](#)

SUMMARY

Senior AI/ML Engineer specialized in production-grade LLM systems and computer vision pipelines. Built high-throughput RAG systems processing tens of thousands of documents daily and scaled infrastructure to support ~40% higher traffic. Experienced in real-time CV models and distributed backend architectures across AWS and Azure.

TECHNICAL SKILLS

LLM / GenAI: RAG, Agentic Systems, LangChain, LlamaIndex, LangGraph, MCP (Model Context Protocol), PEFT, QLoRA, PPO, vLLM, OpenAI, Claude, TTS/STT

CV / ML Frameworks: YOLOv8, OpenCV, PyTorch, TensorFlow, Scikit-learn

Backend: Python, FastAPI, Django, Microservices, REST, GraphQL, n8n, Selenium, BeautifulSoup

Cloud & Data: AWS, Azure, Docker, CI/CD, PostgreSQL, MongoDB, Redis, RabbitMQ

EXPERIENCE

Senior AI/ML Engineer (Part-time) - Devfied (Remote)

Nov 2025 - Present

- Architected a multi-agent RAG system processing 50K+ documents/day with sub-second latency using LangGraph and Azure AI Foundry
- Built production-grade face recognition systems (InsightFace) deployed in real-time verification pipelines, improving accuracy and reducing false positives.
- Automated 8 business workflows using n8n, reducing operational overhead by 15%
- Scaled backend infrastructure using async concurrency models and distributed task queues across AWS/Azure, supporting 40% increase in peak user load.

ML/AI Engineer - Gtek (Remote)

Oct 2025 - Apr 2026

- Led development of multi-agent AI systems, driving 20% increase in user engagement
- Designed real-time CV pipelines (YOLO) achieving 92% accuracy in vehicle tracking
- Improved platform conversion rates by 5% via AI-driven personalization
- Optimized application latency by implementing async concurrency models in the FastAPI backend, reducing response times by 15%.

Senior ML/AI Engineer & Backend Developer - Transdata

Sep 2024 - Sep 2025

- Built scalable microservices (FastAPI, Docker, RabbitMQ), improving CI/CD efficiency
- Designed RAG-based NLP systems processing 10K+ documents/month
- Reduced technical debt by 30% through modular backend architecture
- Improved object detection accuracy by 18% using YOLO-based models
- Maintained 99.9% uptime across multi-cloud infrastructure (AWS, Azure)

ML Engineer & Backend Developer - Software Alliance

Oct 2022 - Oct 2024

- Led AI system roadmap, improving system efficiency by 30% and user satisfaction by 25%
- Developed NLP pipelines and RAG systems for content generation and retrieval
- Fine-tuned LLMs (PEFT, QLoRA, PPO) improving task performance by 20%

EDUCATION

B.Sc. Computer Science - The Superior University, Lahore

CGPA: 3.56 / 4.0

CERTIFICATIONS

Generative AI (Coursera) • Deep Learning (NVIDIA) • Prompt Engineering • Machine Learning with Python